

Jose Tenorio
Coin vs Scrap
EE 5353 Neural Networks

Introduction:

The purpose of this program is to utilize the free cloud service Google Colab. This service allows the user to write and run executable documents. Google Colab connects the user's notebook to a cloud base runtime and execute python code without any setup on the user's machine.

This program executes python code on **Google Colab** that identifies images using **convolutional neural nets** (CNN). CNN's can capture spatial and temporal dependencies of an image through the application of filters. Convolutional neural nets reduce the image into a form which is easier to process without losing features which are critical for a good prediction. This process is done with the use of the **Kernel filter**. This filter is also able to extract high-level features such as edges from the input image. The **pooling layer** is responsible for reducing the spatial size of the convolved feature. This helps decrease the computational power required to process the data. This layer is also useful for extracting dominant features which are rotational and positional invariant. Lastly, the **dropout layer** refers to ignoring units during the training phase of certain neurons which are selected at random. By ignoring, I mean they are not considered during a particular forward or backward pass. We do this to prevent over-fitting.

Procedure:

1. Input the images from the training folder in proper image and label format (use onehot encoding/to_categorical)
2. Divide the data into training (80%) and validation (20%).
3. Resize the images to 200 by 200.
4. Convert the images to black and white
5. Normalize the input data
6. Design a convolutional neural network with the following features:
 - Convolutional layer with 64 filters, Size of the filters is 3, 3, Strides is 2 and relu activation
 - Pooling layer with pool size 2,2
 - Dropout layer with rate 0.5
 - Flattening
 - Dense layer fully connected with 128 hidden units and relu activations
 - Dropout layer with rate 0.5
 - Final dense fully connected layer with number of classes and softmax activation.
7. Verify if the number of iterations/nb_epochs = 20.
8. Validation data should be used during training
9. Input the images from the testing folder in proper image and label format as used for training. Shuffle the testing data.
10. Test the images. Print results and testing figure.

Part I CNN without Augmentation

Training Results:

```
#####  
Load Datasets
```

```
Data Directory List 1- > ['COIN', 'SCRAP']  
Data Directory List 2- > ['COIN', 'SCRAP']  
(array([0, 1]), array([400, 921]))
```

```
#####  
Create Model  
Model: "sequential_4"
```


| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|----------|
| conv2d_4 (Conv2D) | (None, 99, 99, 64) | 640 |
| max_pooling2d_4 (MaxPooling2D) | (None, 49, 49, 64) | 0 |
| dropout_8 (Dropout) | (None, 49, 49, 64) | 0 |
| flatten_4 (Flatten) | (None, 153664) | 0 |
| dense_8 (Dense) | (None, 128) | 19669120 |
| dropout_9 (Dropout) | (None, 128) | 0 |
| dense_9 (Dense) | (None, 2) | 258 |


```
#####  
Total params: 19,670,018  
Trainable params: 19,670,018  
Non-trainable params: 0  
#####
```

```
#####  
Train Model
```

```
Epoch 1/20  
33/33 [=====] - 16s 448ms/step - loss: 1.4788 - accuracy: 0.6029 -  
val_loss: 0.4264 - val_accuracy: 0.8075  
Epoch 2/20  
33/33 [=====] - 14s 430ms/step - loss: 0.4046 - accuracy: 0.8288 -  
val_loss: 0.3578 - val_accuracy: 0.8491  
Epoch 3/20  
33/33 [=====] - 14s 438ms/step - loss: 0.3740 - accuracy: 0.8578 -  
val_loss: 0.3194 - val_accuracy: 0.8642  
Epoch 4/20  
33/33 [=====] - 15s 452ms/step - loss: 0.3435 - accuracy: 0.8636 -  
val_loss: 0.3092 - val_accuracy: 0.8679  
Epoch 5/20  
33/33 [=====] - 14s 438ms/step - loss: 0.3417 - accuracy: 0.8615 -  
val_loss: 0.2774 - val_accuracy: 0.8679  
Epoch 6/20  
33/33 [=====] - 15s 447ms/step - loss: 0.3071 - accuracy: 0.8672 -  
val_loss: 0.2754 - val_accuracy: 0.8717  
Epoch 7/20  
33/33 [=====] - 14s 439ms/step - loss: 0.2847 - accuracy: 0.8779 -  
val_loss: 0.2620 - val_accuracy: 0.8679  
Epoch 8/20  
33/33 [=====] - 14s 435ms/step - loss: 0.2737 - accuracy: 0.8749 -  
val_loss: 0.2500 - val_accuracy: 0.8717  
Epoch 9/20  
33/33 [=====] - 14s 436ms/step - loss: 0.2663 - accuracy: 0.8779 -  
val_loss: 0.2408 - val_accuracy: 0.8792  
Epoch 10/20  
33/33 [=====] - 14s 435ms/step - loss: 0.2478 - accuracy: 0.8874 -  
val_loss: 0.2326 - val_accuracy: 0.8755
```

Epoch 11/20
 33/33 [=====] - 14s 435ms/step - loss: 0.2267 - accuracy: 0.8923 -
 val_loss: 0.2150 - val_accuracy: 0.8830
 Epoch 12/20
 33/33 [=====] - 14s 438ms/step - loss: 0.2133 - accuracy: 0.8977 -
 val_loss: 0.2045 - val_accuracy: 0.9019
 Epoch 13/20
 33/33 [=====] - 14s 438ms/step - loss: 0.2050 - accuracy: 0.9203 -
 val_loss: 0.2040 - val_accuracy: 0.8868
 Epoch 14/20
 33/33 [=====] - 14s 430ms/step - loss: 0.2036 - accuracy: 0.9106 -
 val_loss: 0.1870 - val_accuracy: 0.8981
 Epoch 15/20
 33/33 [=====] - 14s 434ms/step - loss: 0.1819 - accuracy: 0.9316 -
 val_loss: 0.1695 - val_accuracy: 0.9132
 Epoch 16/20
 33/33 [=====] - 14s 432ms/step - loss: 0.1522 - accuracy: 0.9426 -
 val_loss: 0.1652 - val_accuracy: 0.9132
 Epoch 17/20
 33/33 [=====] - 14s 438ms/step - loss: 0.1501 - accuracy: 0.9440 -
 val_loss: 0.1660 - val_accuracy: 0.9094
 Epoch 18/20
 33/33 [=====] - 14s 429ms/step - loss: 0.1503 - accuracy: 0.9425 -
 val_loss: 0.1612 - val_accuracy: 0.9170
 Epoch 19/20
 33/33 [=====] - 15s 442ms/step - loss: 0.1376 - accuracy: 0.9507 -
 val_loss: 0.1482 - val_accuracy: 0.9170
 Epoch 20/20
 33/33 [=====] - 14s 432ms/step - loss: 0.1236 - accuracy: 0.9587 -
 val_loss: 0.1437 - val_accuracy: 0.9170






Testing Results:


```
#####
Test model with validation Data
```

```
9/9 [=====] - 1s 82ms/step - loss: 0.1437 - accuracy: 0.9170
Validation accuracy - > 91.69811606407166
```


The predicted validation image is = SCRAP verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = SCRAP verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below



The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




```
#####  
Test model with test Data
```

```
13/13 [=====] - 1s 93ms/step - loss: 0.1053 - accuracy: 0.9470  
Testing accuracy - > 94.69696879386902
```


The predicted testing image is = COIN verify below




The predicted testing image is = COIN verify below




The predicted testing image is = COIN verify below




The predicted testing image is = SCRAP verify below




The predicted testing image is = SCRAP verify below




The predicted testing image is = COIN verify below




The predicted testing image is = COIN verify below




The predicted testing image is = SCRAP verify below



The predicted testing image is = COIN verify below



The predicted testing image is = COIN verify below



Part II CNN with Augmentation

Training Results:

```
#####  
Load Datasets  
  
Data Directory List 1- > ['COIN', 'SCRAP']  
Data Directory List 2- > ['COIN', 'SCRAP']
```

```
#####  
Create Model
```


```
Model: "sequential_1"
```


| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|----------|
| conv2d_1 (Conv2D) | (None, 99, 99, 64) | 640 |
| max_pooling2d_1 (MaxPooling2D) | (None, 49, 49, 64) | 0 |
| dropout_2 (Dropout) | (None, 49, 49, 64) | 0 |
| flatten_1 (Flatten) | (None, 153664) | 0 |
| dense_2 (Dense) | (None, 128) | 19669120 |
| dropout_3 (Dropout) | (None, 128) | 0 |
| dense_3 (Dense) | (None, 2) | 258 |
| Total params: 19,670,018 | | |
| Trainable params: 19,670,018 | | |
| Non-trainable params: 0 | | |

```
#####  
Train Model
```

```
Epoch 1/20  
33/33 [=====] - 20s 574ms/step - loss: 1.1797 - accuracy: 0.6460 -  
val_loss: 0.4404 - val_accuracy: 0.8415  
Epoch 2/20  
33/33 [=====] - 17s 518ms/step - loss: 0.4447 - accuracy: 0.8299 -  
val_loss: 0.3561 - val_accuracy: 0.8415  
Epoch 3/20  
33/33 [=====] - 15s 462ms/step - loss: 0.3565 - accuracy: 0.8516 -  
val_loss: 0.3273 - val_accuracy: 0.8679  
Epoch 4/20  
33/33 [=====] - 15s 461ms/step - loss: 0.3574 - accuracy: 0.8505 -  
val_loss: 0.2944 - val_accuracy: 0.8717  
Epoch 5/20  
33/33 [=====] - 15s 458ms/step - loss: 0.3371 - accuracy: 0.8717 -  
val_loss: 0.2926 - val_accuracy: 0.8717  
Epoch 6/20  
33/33 [=====] - 15s 453ms/step - loss: 0.3063 - accuracy: 0.8817 -  
val_loss: 0.2782 - val_accuracy: 0.8830  
Epoch 7/20  
33/33 [=====] - 15s 462ms/step - loss: 0.3447 - accuracy: 0.8573 -  
val_loss: 0.2727 - val_accuracy: 0.8717  
Epoch 8/20  
33/33 [=====] - 15s 459ms/step - loss: 0.3386 - accuracy: 0.8678 -  
val_loss: 0.2537 - val_accuracy: 0.8792  
Epoch 9/20  
33/33 [=====] - 15s 460ms/step - loss: 0.2848 - accuracy: 0.8903 -  
val_loss: 0.2557 - val_accuracy: 0.8642  
Epoch 10/20  
33/33 [=====] - 15s 458ms/step - loss: 0.3095 - accuracy: 0.8812 -  
val_loss: 0.2307 - val_accuracy: 0.8981  
Epoch 11/20  
33/33 [=====] - 15s 457ms/step - loss: 0.2868 - accuracy: 0.8815 -  
val_loss: 0.2163 - val_accuracy: 0.9208  
Epoch 12/20  
33/33 [=====] - 15s 453ms/step - loss: 0.2601 - accuracy: 0.8948 -  
val_loss: 0.2076 - val_accuracy: 0.9245  
Epoch 13/20  
33/33 [=====] - 15s 453ms/step - loss: 0.2512 - accuracy: 0.8946 -  
val_loss: 0.2061 - val_accuracy: 0.9094  
Epoch 14/20  
33/33 [=====] - 15s 455ms/step - loss: 0.2129 - accuracy: 0.9063 -  
val_loss: 0.2302 - val_accuracy: 0.9094
```


Epoch 15/20
 33/33 [=====] - 15s 452ms/step - loss: 0.2096 - accuracy: 0.9154 -
 val_loss: 0.1876 - val_accuracy: 0.9321
 Epoch 16/20
 33/33 [=====] - 15s 454ms/step - loss: 0.2228 - accuracy: 0.9232 -
 val_loss: 0.1726 - val_accuracy: 0.9283
 Epoch 17/20
 33/33 [=====] - 15s 453ms/step - loss: 0.2324 - accuracy: 0.9171 -
 val_loss: 0.1541 - val_accuracy: 0.9283
 Epoch 18/20
 33/33 [=====] - 15s 455ms/step - loss: 0.1892 - accuracy: 0.9167 -
 val_loss: 0.1654 - val_accuracy: 0.9472
 Epoch 19/20
 33/33 [=====] - 15s 454ms/step - loss: 0.2087 - accuracy: 0.9187 -
 val_loss: 0.1426 - val_accuracy: 0.9547
 Epoch 20/20
 33/33 [=====] - 15s 450ms/step - loss: 0.1881 - accuracy: 0.9327 -
 val_loss: 0.1507 - val_accuracy: 0.9245






Testing Results:


```
#####
Test model with validation Data
```

```
9/9 [=====] - 1s 88ms/step - loss: 0.1507 - accuracy: 0.9245
Validation accuracy - > 92.45283007621765
```


The predicted validation image is = SCRAP verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = SCRAP verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = COIN verify below




The predicted validation image is = SCRAP verify below



The predicted validation image is = COIN verify below



The predicted validation image is = COIN verify below




```
#####
```

```
Test model with test Data
```


```
13/13 [=====] - 1s 98ms/step - loss: 0.1612 - accuracy: 0.9369
```

```
Testing accuracy - > 93.6868667602539
```


```
The predicted testing image is = COIN verify below
```




```
The predicted testing image is = COIN verify below
```




The predicted testing image is = COIN verify below




The predicted testing image is = SCRAP verify below




The predicted testing image is = COIN verify below




The predicted testing image is = COIN verify below




The predicted testing image is = COIN verify below




The predicted testing image is = SCRAP verify below



The predicted testing image is = COIN verify below



The predicted testing image is = COIN verify below



Input/Output Processing Explanation:

Step 1: Import data

Step 2: Define number of classes

Step 3:

For data set in directory

 For images in image list

 Read image

 Change color space

 Resize

 Append label

 End

End

Step 4: Output unique labels

Step 5: Shuffle dataset

Step 6: Divide data into train & test

Step 7: Normalize data

Step 8: Reshape data to fit model

Step 9: Add convolutionary, max pooling, and dropout layers

Step 10: Compile Model

Step 11: Fit model on training Data

Step 12: Evaluate model on test data

Step 13: Test accuracy

Step 14: Print images vs predicted

Code:

<https://github.com/jnn-dev>

Conclusion:

This programming exercise utilized convolutional neural nets with augmentation to capture spatial and temporal dependencies of images through the application of filters. The augmentation included a horizontal and vertical flip to the training dataset. This augmentation helped improved the validation accuracy, but slightly lowered the training accuracy.